

Uncertainty-Aware Hybrid Rendering with Gaussian Splatting and NeRF for High-Fidelity Synthesis

Team 25 Patrick Chen

Robotics Institute, School of Computer Science, Carnegie Mellon University



Introduction




Figure 1. NeRF VS 3D GS

- Real-time rendering in 3D scenes often trades off between speed and quality.
- 3D Gaussian Splatting (3D GS)** is fast but may lack fine detail.
- NeRF** offers high fidelity but is computationally expensive.
- We propose a **hybrid rendering** pipeline that combines both, guided by **uncertainty estimation**.

Related Work

This project builds upon existing work in NeRF and Gaussian Splatting:

- NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis** [3] – Introduced NeRF for novel view synthesis but suffers from slow inference.
- 3D Gaussian Splatting for Real-Time Radiance Field Rendering** [2] – Enabled real-time rendering but lacks high-resolution details.
- SS Mip-NeRF: Supersampled Mip-NeRF** [1] – Supersampling improves Mip-NeRF accuracy with minimal rendering overhead.

Carnegie Mellon University
School of Computer Science

Method

We propose a hybrid rendering framework that leverages the speed of 3D Gaussian Splatting (GS) and the fidelity of NeRF through uncertainty-guided pixel-wise rendering. The full framework consists of three main stages:




Figure 2. Uncertainty-Aware Hybrid Rendering Framework

- 3D GS Rendering:** Train a 3D GS model with multi-view RGB and poses, then render views via efficient splatting.
- Uncertainty Prediction:** Use a UNet/MLP to predict per-pixel SSIM error from GS features, yielding an uncertainty map.
- Pixel-Wise NeRF Rendering:** Re-render top- k % uncertain pixels with NeRF; merge with GS to form the final output.

Neural Uncertainty Feature

We extract four features from 3D GS: alpha sum, color variance, view direction, and 2D covariance area. These capture visibility, appearance uncertainty, angular sensitivity, and footprint scale, respectively, and are used to predict per-pixel uncertainty.




Figure 3. Feature Extraction for Neural Uncertainty Prediction

Qualitative Result

We visualize hybrid rendering results under a 30% uncertainty-based replacement strategy. The predicted uncertainty highlights regions where GS struggles, allowing NeRF to selectively refine them and enhance perceptual quality efficiently.




Figure 4. Qualitative results with top 30% highest uncertainty replacement.

Quantitative Result

GS time dominates when $k \leq 30\%$, but NeRF dominates for $k > 30\%$. The 1K MLP (Spearman: 0.756) performs similarly to the 110K U-Net (0.890). **Hybrid-MLP 50%** achieves top (PSNR / SSIM) = (24.75 / 0.892) and fastest time (0.668s), outperforming GS-only (22.23 / 0.819), NeRF-only (24.76 / 0.891, 0.827s), and SS Mip-NeRF (0.864s), and is slightly faster than **Hybrid-UNet 50%** (0.673s). This shows smaller models offer efficient, high-quality hybrid rendering.

Method	PSNR↑	SSIM↑	GS Time (s)	NeRF Time (s)	Total Time (s)	Replace %
GS only	22.23	0.819	0.246	—	0.246	0%
NeRF only	24.76	0.891	—	0.827	0.827	100%
SS Mip-NeRF	32.60	—	—	0.864	0.864	100%
Hybrid-MLP 5%	22.39	0.824	0.272	0.040	0.312	~5%
Hybrid-MLP 10%	22.64	0.832	0.272	0.081	0.353	~10%
Hybrid-MLP 30%	24.31	0.875	0.272	0.220	0.492	~30%
Hybrid-MLP 50%	24.75	0.892	0.271	0.397	0.668	~50%
Hybrid-UNet 5%	22.58	0.836	0.275	0.040	0.315	~5%
Hybrid-UNet 10%	22.88	0.847	0.274	0.079	0.353	~10%
Hybrid-UNet 30%	24.11	0.880	0.274	0.220	0.494	~30%
Hybrid-UNet 50%	24.73	0.892	0.274	0.399	0.673	~50%

Table 1. Rendering quality and runtime at different replacement rates using MLP vs. UNet for uncertainty prediction.

Note: GS Time for Hybrid refers to the time spent on feature extraction using GS model. NeRF Time for Hybrid refers to the time for ray-marching rendering using NeRF.

References

- Jonathan T. Barron, Ben Mildenhall, Matthew Tancik, Peter Hedman, Ricardo Martin-Brualla, and Pratul P. Srinivasan. Mip-nerf: A multiscale representation for anti-aliasing neural radiance fields, 2021.
- Bernhard Kerbl, Georgios Kopanas, Thomas Leimkühler, and George Drettakis. 3d gaussian splatting for real-time radiance field rendering. *ACM Transactions on Graphics*, 42(4), July 2023.
- Ben Mildenhall, Pratul P. Srinivasan, Matthew Tancik, Jonathan T. Barron, Ravi Ramamoorthi, and Ren Ng. Nerf: Representing scenes as neural radiance fields for view synthesis. In *ECCV*, 2020.